

Member	ID	Badge
محمد مسعد نعيم محمد	23011492	Linear Regression
عبدالرحمن احمد عبدالمنعم	23011311	Polynomial Regression
محمد علاء محمد احمد البرعي	23011481	Logistic Regression

For The Source Code - Click Here for Github Repo

Import libraries

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.decomposition import PCA
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.preprocessing import StandardScaler, LabelEncoder, PolynomialFeatures
from sklearn.metrics import mean_squared_error, r2_score, confusion_matrix, Confusi
```

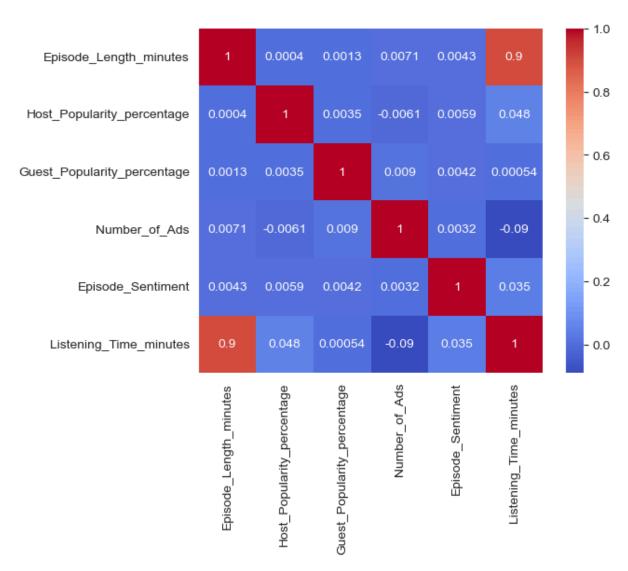
Exploratory Data Analysis (EDA)

```
In [5]: SM_data = pd.read_csv("https://raw.githubusercontent.com/mohamed8905/Regression-Mod
SM_data
```

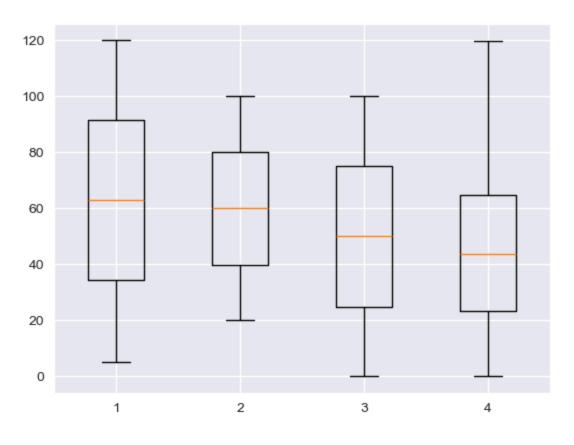
```
Out[5]:
                Podcast_Name Episode_Title Episode_Length_minutes
                                                                      Genre Host_Popularity_per
             0
                 Healthy Living
                                  Episode 77
                                                              99.25
                                                                      Health
                                                                       True
                      Mystery
             1
                                  Episode 6
                                                              19.43
                       Matters
                                                                      Crime
                 Current Affairs
                                  Episode 1
                                                             117.03
                                                                       News
                                                                       True
                      Mystery
             3
                                  Episode 38
                                                              16.97
                       Matters
                                                                      Crime
             4
                   Humor Hub
                                 Episode 73
                                                              83.48 Comedy
         52495
                Home & Living
                                 Episode 17
                                                              24.81
                                                                    Lifestyle
         52496
                   Melody Mix
                                  Episode 9
                                                              92.15
                                                                      Music
                      Comedy
         52497
                                  Episode 24
                                                             112.27 Comedy
                       Corner
         52498
                 Business Briefs
                                 Episode 85
                                                              NaN Business
         52499
                 Healthy Living
                                 Episode 54
                                                              94.35
                                                                      Health
        52500 rows × 11 columns
In [6]: SM data['Episode Sentiment'] = SM data['Episode Sentiment'].map({'Negative':-1, 'Ne
        SM_data.info()
       <class 'pandas.core.frame.DataFrame'>
       RangeIndex: 52500 entries, 0 to 52499
       Data columns (total 11 columns):
           Column
                                          Non-Null Count Dtype
       ---
           _____
        0
            Podcast_Name
                                          52500 non-null object
            Episode Title
                                          52500 non-null
                                                           object
        2
            Episode_Length_minutes
                                          47254 non-null float64
        3
                                          52500 non-null
                                                           object
            Genre
                                          52500 non-null float64
            Host_Popularity_percentage
        5
            Publication_Day
                                          52500 non-null
                                                           object
            Publication_Time
                                          52500 non-null
                                                           object
        7
            Guest_Popularity_percentage 47242 non-null
                                                           float64
            Number_of_Ads
                                          52500 non-null
                                                           int64
            Episode_Sentiment
                                          52500 non-null
                                                           int64
        10 Listening_Time_minutes
                                          47105 non-null float64
       dtypes: float64(4), int64(2), object(5)
       memory usage: 4.4+ MB
In [7]: # Dropping null values
         null_rows = SM_data.isnull().any(axis=1).sum()
         print(f"Rows with at least one null: {null_rows}")
        SM_data[SM_data.isnull().any(axis=1)]
```

Rows with at least one null: 14337

```
Out[7]:
                Podcast_Name Episode_Title Episode_Length_minutes
                                                                        Genre Host_Popularity_I
             4
                   Humor Hub
                                 Episode 73
                                                             83.48
                                                                      Comedy
                Money Matters
                                 Episode 87
                                                             28.06
                                                                      Business
            10
                  Joke Junction
                                 Episode 68
                                                             34.97
                                                                      Comedy
            22
                                                                   Technology
                  Digital Digest
                                 Episode 48
                                                             85.64
            23
                  Learning Lab
                                 Episode 18
                                                              NaN
                                                                     Education
         52474 News Roundup
                                 Episode 66
                                                             44.52
                                                                        News
         52476
                   Health Hour
                                 Episode 16
                                                             99.19
                                                                        Health
        52491
                    Sport Spot
                                 Episode 86
                                                             29.47
                                                                        Sports
         52493
                   Brain Boost
                                 Episode 71
                                                                     Education
                                                             72.87
        52498
                 Business Briefs
                                 Episode 85
                                                              NaN
                                                                      Business
        14337 rows × 11 columns
        SM_data.dropna(inplace=True, axis=0)
In [8]:
        SM_data.info()
       <class 'pandas.core.frame.DataFrame'>
       Index: 38163 entries, 0 to 52499
       Data columns (total 11 columns):
            Column
        #
                                          Non-Null Count Dtype
       _ _ _
            -----
                                          -----
            Podcast_Name
                                          38163 non-null object
        0
        1
            Episode_Title
                                          38163 non-null object
                                          38163 non-null float64
        2
            Episode_Length_minutes
        3
            Genre
                                          38163 non-null object
            Host_Popularity_percentage
                                          38163 non-null float64
        5
            Publication_Day
                                          38163 non-null
                                                          object
                                                          object
        6
            Publication_Time
                                          38163 non-null
        7
            Guest_Popularity_percentage 38163 non-null
                                                          float64
            Number_of_Ads
                                          38163 non-null
                                                          int64
            Episode_Sentiment
                                          38163 non-null
                                                          int64
        10 Listening_Time_minutes
                                          38163 non-null float64
       dtypes: float64(4), int64(2), object(5)
       memory usage: 3.5+ MB
In [9]:
        # Heatmap to find correlation
        num_DM_data = SM_data.select_dtypes(include=["number"])
        sns.heatmap(num_DM_data.corr(), annot=True, square=True, cmap='coolwarm')
        plt.show()
```



```
In [10]: # Checking outliers
plt.boxplot(num_DM_data.drop(['Number_of_Ads', 'Episode_Sentiment'], axis=1))
plt.show()
```



Preprocessing

```
In [11]: # Splitting data to input / output
x = num_DM_data.drop('Listening_Time_minutes', axis=1)
y = num_DM_data['Listening_Time_minutes']
x
```

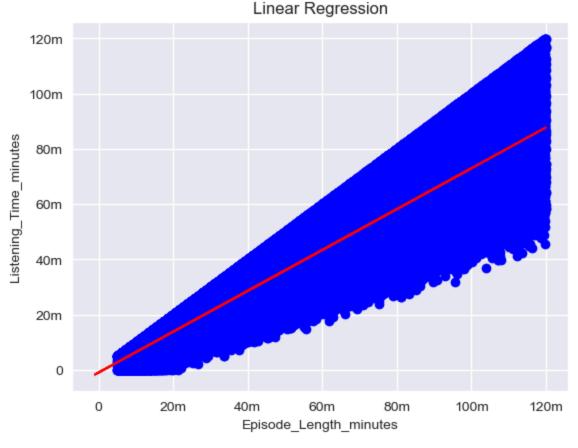
Out[11]:		Episode_Length_minutes	Host_Popularity_percentage	Guest_Popularity_percentage I
	0	99.25	21.37	70.22
	1	19.43	47.19	75.15
	2	117.03	96.33	57.95
	3	16.97	25.73	24.19
	6	63.77	58.84	82.20
	•••			
	52494	33.35	63.23	90.81
	52495	24.81	66.15	98.63
	52496	92.15	89.61	25.82
	52497	112.27	26.33	55.29
	52499	94.35	91.08	92.54
	38163 rd	ows × 5 columns		
	4			-
In [12]:	у			

```
Out[12]: 0
                 55.158695
                   7.686559
                  110.064645
         3
                  12.000380
                   57.827346
         52494 32.133509
         52495
                  20.573795
         52496
                  76.198459
         52497
                  107.602135
         52499
                   52.102024
         Name: Listening_Time_minutes, Length: 38163, dtype: float64
In [13]: # Splitting data to train / test
         x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2)
```

Linear Regression (Simple)

```
In [14]: # Building Simple Linear model
    reg = LinearRegression()
    reg.fit(x_train[["Episode_Length_minutes"]], y_train)
    print(f"The Equation is: Y = {reg.coef_[0]:0.3f} X + {reg.intercept_:0.3f}")
The Equation is: Y = 0.739 X + -0.905
```

```
In [15]: plt.scatter(x_train[["Episode_Length_minutes"]], y_train , color='blue')
    plt.plot(x_train, reg.coef_[0]*x_train + reg.intercept_, '-r')
    plt.xlabel('Episode_Length_minutes')
    plt.ylabel('Listening_Time_minutes')
    plt.title("Linear Regression")
    plt.xticks([0,20,40,60,80,100,120], ['0','20m','40m','60m','80m','100m','120m'])
    plt.yticks([0,20,40,60,80,100,120], ['0','20m','40m','60m','80m','100m','120m'])
    plt.show()
```



```
In [16]: # Making prediction
    y_hat = reg.predict(x_test[["Episode_Length_minutes"]])
In [17]: # testing score and error
    r2 = r2_score(y_test, y_hat)
    print(f"Model Accuracy (R² Score): {r2*100:.2f}%")

    Model Accuracy (R² Score): 81.20%
In [18]: mse = mean_squared_error(y_test, y_hat)
    print(f"Mean Squared Error: {mse:.2f}")

    Mean Squared Error: 137.99
```

Multiple Linear Regression

```
In [19]: # Building Multiple Linear model
          reg2 = LinearRegression()
          reg2.fit(x_train, y_train)
          for i in range(len(reg2.coef_)):
              if i==0:
                  eq = eq + f"Y = {reg2.intercept_:0.3f} + {reg2.coef_[i]:0.3f} X{i+1}"
                  eq = eq + f'' + \{reg2.coef_[i]:0.3f\} X\{i+1\}''
          print(f"The Equation is: {eq}")
        The Equation is: Y = -0.790 + 0.739 \times 1 + 0.056 \times 2 + 0.000 \times 3 + -2.343 \times 4 + 1.017 \times 5
In [20]: # Prediction
         y_hat2 = reg2.predict(x_test)
In [21]: # Score and error
          r2_2 = r2_score(y_test, y_hat2)
          print(f"Model Accuracy (R2 Score): {r2_2*100:.2f}%")
        Model Accuracy (R2 Score): 82.41%
In [22]: mse2 = mean_squared_error(y_test, y_hat)
          print(f"Mean Squared Error: {mse2:.2f}")
        Mean Squared Error: 137.99
```

The Multiple Linear Regression Model has higher r2 score.

Polynomial Regression

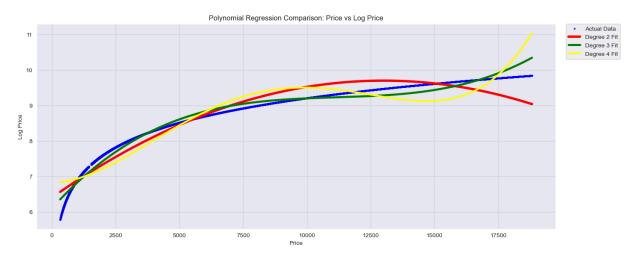
```
In [23]: Poly_reg = pd.read_csv("https://raw.githubusercontent.com/mohamed8905/Regression-Mo
         # Clean the data: Drop missing or duplicate rows
         Poly_reg = Poly_reg.dropna().drop_duplicates()
         # Convert columns to numeric
         Poly_reg["price"] = pd.to_numeric(Poly_reg["price"], errors='coerce')
         Poly_reg["log_price"] = pd.to_numeric(Poly_reg["log_price"], errors='coerce')
         Poly_reg = Poly_reg.dropna(subset=["price", "log_price"])
         Poly_reg
```

Out[23]:		carat	cut	color	clarity	depth	table	price	X	у	z	log_price	1
	0	0.23	Ideal	Е	SI2	61.5	55.0	326	3.95	3.98	2.43	5.786897	38.
	1	0.21	Premium	Е	SI1	59.8	61.0	326	3.89	3.84	2.31	5.786897	34.
	2	0.23	Good	Е	VS1	56.9	65.0	327	4.05	4.07	2.31	5.789960	38.
	3	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63	5.811141	46.
	4	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	2.75	5.814131	51.
	•••												
	53935	0.72	Ideal	D	SI1	60.8	57.0	2757	5.75	5.76	3.50	7.921898	115.
	53936	0.72	Good	D	SI1	63.1	55.0	2757	5.69	5.75	3.61	7.921898	118.
	53937	0.70	Very Good	D	SI1	62.8	60.0	2757	5.66	5.68	3.56	7.921898	114.
	53938	0.86	Premium	Н	SI2	61.0	58.0	2757	6.15	6.12	3.74	7.921898	140.
	53939	0.75	Ideal	D	SI2	62.2	55.0	2757	5.83	5.87	3.64	7.921898	124.

53794 rows × 12 columns

```
In [24]: # Comparing 2nd, 3rd & 4th degree
         X = Poly_reg["price"].values.reshape(-1, 1)
         y = Poly_reg["log_price"].values
         degrees = [2, 3, 4]
         models = {}
         predictions = {}
         metrics = {}
         for degree in degrees:
             poly = PolynomialFeatures(degree=degree)
             X_poly = poly.fit_transform(X)
             model = LinearRegression()
             model.fit(X_poly, y)
             y_pred = model.predict(X_poly)
             models[degree] = model
             predictions[degree] = y_pred
             mse = mean_squared_error(y, y_pred)
             r2 = r2\_score(y, y\_pred)
             metrics[degree] = {
                  'MSE': mse,
                  'R2': r2,
                  'Accuracy%': r2 * 100
             }
```

```
In [25]: for degree in degrees:
              print(f"\nDegree {degree} Results:")
              print(f"Mean Squared Error: {metrics[degree]['MSE']:.4f}")
              print(f"R2 Score: {metrics[degree]['R2']:.4f}")
              print(f"Accuracy: {metrics[degree]['Accuracy%']:.2f}%")
        Degree 2 Results:
        Mean Squared Error: 0.0535
        R<sup>2</sup> Score: 0.9480
        Accuracy: 94.80%
        Degree 3 Results:
        Mean Squared Error: 0.0181
        R<sup>2</sup> Score: 0.9824
        Accuracy: 98.24%
        Degree 4 Results:
        Mean Squared Error: 0.1106
        R<sup>2</sup> Score: 0.8925
        Accuracy: 89.25%
In [26]: plt.figure(figsize=(13.5, 6))
         plt.scatter(X.flatten(), y, color='blue', s=6, alpha=0.6, label='Actual Data')
         # Plot polynomial fits
          colors = ['red', 'green', 'yellow']
          sorted_idx = np.argsort(X.flatten())
         for i, degree in enumerate(degrees):
              plt.plot(X[sorted_idx].flatten(), predictions[degree][sorted_idx],
                       color=colors[i], linewidth=4-0.5*i,
                       label=f'Degree {degree} Fit')
          plt.title('Polynomial Regression Comparison: Price vs Log Price')
          plt.xlabel('Price')
         plt.ylabel('Log Price')
          # Adjust Legend position
          plt.legend(bbox_to_anchor=(1.02, 1), loc='upper left', borderaxespad=0.)
          # Grid settings
          plt.grid(True, color='LightGray', linewidth=1)
          # Adjust margins
          plt.subplots_adjust(left=0.05, right=0.95, bottom=0.1, top=0.9)
          plt.show()
```



```
In [27]: # The models
for degree in degrees:
    print(f"\nDegree {degree} Polynomial Equation:")
    equation = f"Y = {models[degree].intercept_}"
    for power, coef in enumerate(models[degree].coef_[1:], 1):
        equation += f" + {coef} X^{power}"
    print(equation)
```

Degree 2 Polynomial Equation:

 $Y = 6.404746158769985 + 0.0005069698301360856 X^1 + -1.9492525316013317e-08 X^2$

Degree 3 Polynomial Equation:

 $Y = 6.103030973684653 + 0.0007968307881509356 X^1 + -6.94722237262789e-08 X^2 + 2.078127513751212e-12 X^3$

Degree 4 Polynomial Equation:

 $Y = 6.830825765370128 + 3.155340445445842e-11 X^1 + 1.2340557074382692e-07 X^2 + -1.3904774232973876e-11 X^3 + 4.239562356902578e-16 X^4$

Logistic Regression with Classification

In [28]: Log_reg=pd.read_csv("https://raw.githubusercontent.com/mohamed8905/Regression-Model
Log_reg

Out[28]:		Booking_ID	no_of_adults	no_of_children	no_of_weeken	d_nights	no_of_week_nights		
	0	INN00001	2	0		1	2		
	1	INN00002	2	0		2	3		
	2	INN00003	1	0		2	1		
	3	INN00004	2	0		0	2		
	4	INN00005	2	0		1	1		
	36270	INN36271				2			
	36271	INN36271	2	0		1	3		
	36272	INN36273	2	0		2	6		
	36273	INN36274	2	0		0	3		
	36274	INN36275	2	0		1	2		
	36275 rd	ows × 19 colu	mns						
	4						•		
In [29]:	log re	g info()							
	RangeInd # Col # Col 0 Bood 1 no 2 no 3 no 4 no 5 typ 6 red 7 rod 8 lea 9 arr 10 arr 11 arr 12 mar 13 rep 14 no 15 no 16 avg 17 no 18 bood dtypes:	dex: 36275 er Lumns (total Lumn of_adults of_children of_weekend_r of_week_nigh oe_of_meal_p duired_car_pa om_type_reser ad_time rival_year rival_month rival_date rket_segment oeated_guest of_previous of_previous of_precial_r oking_status	nts lan arking_space rved _type _cancellation _bookings_not room requests int64(13), o	36274 Non 362 362 362 362 362 362 362 362 362 362	75 non-null : 75				
	<pre>In [30]: Log_reg.describe()</pre>								
_ =		-							

```
Out[30]:
                 no_of_adults no_of_children no_of_weekend_nights no_of_week_nights required_ca
          count 36275.000000
                                36275.000000
                                                      36275.000000
                                                                         36275.000000
                     1.844962
                                                          0.810724
          mean
                                    0.105279
                                                                             2.204300
            std
                     0.518715
                                    0.402648
                                                          0.870644
                                                                             1.410905
            min
                     0.000000
                                    0.000000
                                                          0.000000
                                                                             0.000000
           25%
                     2.000000
                                    0.000000
                                                          0.000000
                                                                             1.000000
           50%
                     2.000000
                                    0.000000
                                                          1.000000
                                                                             2.000000
           75%
                     2.000000
                                    0.000000
                                                          2.000000
                                                                             3.000000
           max
                     4.000000
                                   10.000000
                                                          7.000000
                                                                            17.000000
In [31]:
          Log reg.isna().sum()
Out[31]:
          Booking_ID
                                                    0
                                                    0
          no_of_adults
          no_of_children
                                                    0
          no_of_weekend_nights
                                                    0
          no_of_week_nights
                                                    0
          type_of_meal_plan
                                                    0
          required_car_parking_space
                                                    0
          room_type_reserved
                                                    0
          lead time
                                                    0
          arrival year
                                                    0
          arrival_month
                                                    0
          arrival date
                                                    0
          market_segment_type
                                                    0
          repeated_guest
                                                    0
          no_of_previous_cancellations
                                                    0
          no_of_previous_bookings_not_canceled
                                                    0
                                                    0
          avg_price_per_room
                                                    0
          no_of_special_requests
          booking_status
                                                    0
          dtype: int64
In [32]: Log_reg.duplicated().sum()
Out[32]: np.int64(0)
In [33]: # Encoding
          #label encoding to type_of_meal_plan
          Log_reg['type_of_meal_plan']=Log_reg['type_of_meal_plan'].str.replace("Meal Plan","
          Log_reg['type_of_meal_plan']=Log_reg['type_of_meal_plan'].str.replace("Not Selected
          Log_reg['type_of_meal_plan']=Log_reg['type_of_meal_plan'].astype(int)
          #label encoding to room_type_reserved
          Log_reg['room_type_reserved']=Log_reg['room_type_reserved'].str.replace("Room_Type"
          Log_reg['room_type_reserved']=Log_reg['room_type_reserved'].astype(int)
          #label encoding to booking_status
          encoder = LabelEncoder()
```

```
Log_reg['booking_status'] = encoder.fit_transform(Log_reg['booking_status'])
Log_reg['booking_status']=Log_reg['booking_status'].astype(int)
Log_reg['booking_status']=Log_reg['booking_status']^1
#label encoding to market_segment_type
encoder = LabelEncoder()
Log_reg['market_segment_type'] = encoder.fit_transform(Log_reg['market_segment_type']
Log_reg['booking_status']=Log_reg['booking_status'].astype(int)
Log_reg
```

Out[33]:

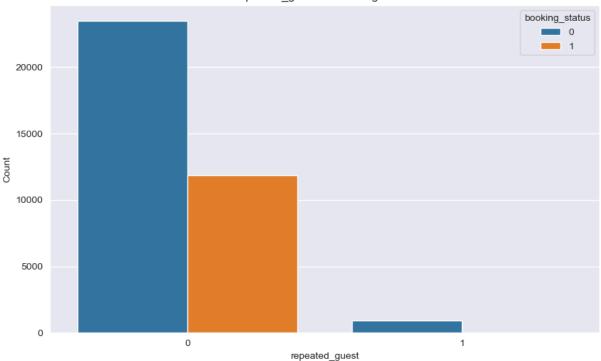
,		Booking_ID	no_of_adults	no_of_children	no_of_weekend_nights	no_of_week_nights
	0	INN00001	2	0	1	2
	1	INN00002	2	0	2	3
	2	INN00003	1	0	2	1
	3	INN00004	2	0	0	2
	4	INN00005	2	0	1	1
	•••	•••				
	36270	INN36271	3	0	2	6
	36271	INN36272	2	0	1	3
	36272	INN36273	2	0	2	6
	36273	INN36274	2	0	0	3
	36274	INN36275	2	0	1	2

36275 rows × 19 columns

```
In [34]: # Feature Engineering
         Log_reg['num_of_days_spend']=Log_reg['no_of_weekend_nights']+Log_reg['no_of_week_ni
         Log_reg['total_peaple']=Log_reg['no_of_adults']+Log_reg['no_of_children']
```

```
In [35]: plt.figure(figsize=(10, 6))
         sns.countplot(data=Log_reg, x='repeated_guest', hue='booking_status')
         plt.title('repeated_guest vs Booking Status')
         plt.xlabel('repeated_guest')
         plt.ylabel('Count')
         plt.show()
```



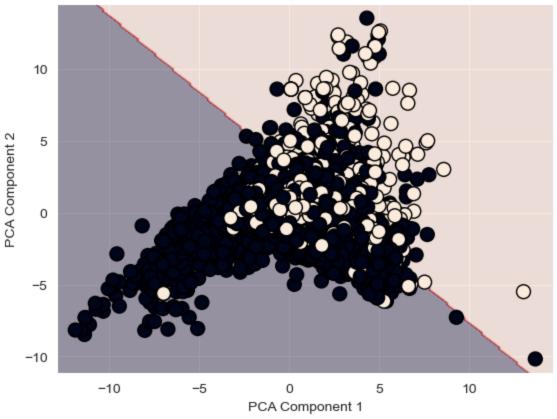


```
In [36]: # Heatmap to show correlation between features
    normacalCol=Log_reg.drop(['Booking_ID'],axis=1)
    plt.figure(figsize=(10, 7))
    sns.heatmap(
        normacalCol.corr(), annot=True, vmin=-1, vmax=1, fmt=".2f", cmap="Spectral"
    )
    plt.show()
```

```
no_of_adults 1.00-0.02 0.10 0.11 -0.00 0.01 0.27 0.10 0.08 0.02 0.03 0.31 -0.19-0.05-0.12 0.30 0.19 0.09 0.13 0.79
                              no_of_children -0.02 1.00 0.03 0.02 0.04 0.03 0.36 -0.05 0.05 -0.00 0.03 0.13 -0.04-0.02-0.02 0.34 0.12 0.03 0.03 0.60
                       - 0.75
                          type_of_meal_plan -0.00 0.04 -0.02 0.03 1.00 -0.02 0.09 0.23 -<mark>0.19</mark> 0.02 0.02 -<mark>0.25</mark> 0.01 -0.01 0.01 0.13 -<mark>0.09</mark> 0.05 0.01 0.02
                                                                                                                           0.50
                    required_car_parking_space 0.01 0.03 -0.03 -0.05 -0.02 1.00 0.04 -0.07 0.02 -0.02 -0.00 -0.00 0.11 0.03 0.06 0.06 0.09 -0.09 -0.05 0.03
                         room_type_reserved 0.27 0.36 0.06 0.09 0.09 0.04 1.00 0.11 0.10 0.01 0.03 0.16 0.03 0.01 0.01 0.47 0.15 0.02 0.10 0.44
                                 lead_time 0.10-0.05 0.05 0.15 0.23-0.07-0.11 1.00 0.14 0.14 0.01-0.01-0.14-0.05-0.08-0.06-0.10 0.44 0.14 0.05
                                                                                                                           - 0.25
                                arrival_year 0.08 0.05 0.06 0.03 -0.19 0.02 0.10 0.14 1.00 -0.34 0.02 0.15 -0.02 0.00 0.03 0.18 0.05 0.18 0.05 0.09
                              arrival_month 0.02-0.00-0.01 0.04 0.02-0.02-0.01 0.14 0.34 1.00 0.04-0.01 0.00 -0.04-0.01 0.05 0.11 -0.01 0.02 0.02
                                                                                                                           - 0.00
                                arrival_date 0.03 0.03 0.03 -0.01 0.02 -0.00 0.03 0.01 0.02 -0.04 1.00 0.01 -0.02-0.01-0.00 0.02 0.02 0.01 0.01 0.04
                        repeated_guest -0.19-0.04-0.07-0.10 0.01 0.11 -0.03-0.14-0.02 0.00 -0.02-0.34 1.00 0.39 0.54 -0.17-0.01-0.11-0.11-0.18
                                                                                                                           - -0.25
                   no_of_previous_bookings_not_canceled -0.12-0.02-0.03-0.05 0.01 0.06-0.01-0.08 0.03-0.01-0.00-0.21 0.54 0.47 1.00 -0.11 0.03-0.06-0.05-0.11
                                                                                                                           - -0.50
                         avg_price_per_room 0.30 0.34 -0.00 0.02 0.13 0.06 0.47 -0.06 0.18 0.05 0.02 0.38 -0.17-0.06-0.11 1.00 0.18 0.14 0.02 0.45
                       no_of_special_requests 0.19 0.12 0.06 0.05 -0.09 0.09 0.15 -0.10 0.05 0.11 0.02 0.31 -0.01-0.00 0.03 0.18 1.00 -0.25 0.07 0.23
                             booking_status 0.09 0.03 0.06 0.09 0.05 -0.09 0.02 0.44 0.18 -0.01 0.01 0.14 -0.11 -0.03 -0.06 0.14 -0.25 1.00 0.10 0.09
                                                                                                                            -0.75
                         num_of_days_spend 0.13 0.03 0.63 0.88 0.01 0.05 0.10 0.14 0.05 0.02 0.01 0.15 0.11 0.03 0.05 0.02 0.07 0.10 1.00 0.13
                               total_peaple 0.79 0.60 0.10 0.10 0.02 0.03 0.44 0.05 0.09 0.02 0.04 0.33 0.18 0.05 0.11 0.45 0.23 0.09 0.13 1.00
                                                                                                                            -1 00
                                                                               arrival_dat
                                                                                   market_segment_typ
                                                     no_of_week_night
                                                                                                     no of special request
                                                  no_of_weekend_nigh
                                                             equired_car_parking_spac
                                                                                          no_of_previous_cancellati
                                                                                              previous_bookings_not_cance
In [37]: # Split data to train and test
             x_data=normacalCol.drop(['booking_status'],axis=1)
             y_data=Log_reg['booking_status']
             x_train, x_test, y_train, y_test = train_test_split(x_data,y_data, test_size=0.2, r
In [38]: # Feature Scaling
             scaler = StandardScaler()
             x_train=scaler.fit_transform(x_train)
             x_test=scaler.transform(x_test)
In [39]: # Fit data to make Logistic Regression
             log=LogisticRegression()
             log.fit(x_train,y_train)
             # Prediction
             predict=log.predict(x_test)
In [40]: # Logistic Regression Decision Boundary with PCA
             pca = PCA(n_components=2)
             X_pca = pca.fit_transform(x_train)
             xx, yy = np.meshgrid(np.linspace(X_pca[:, 0].min() - 1, X_pca[:, 0].max() + 1, 200)
                                          np.linspace(X_pca[:, 1].min() - 1, X_pca[:, 1].max() + 1, 200)
             Z = log.predict(pca.inverse_transform(np.c_[xx.ravel(), yy.ravel()]))
             Z = Z.reshape(xx.shape)
             plt.contourf(xx, yy, Z, alpha=0.4)
             plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y_train, edgecolors='k', marker='o', s=100)
```

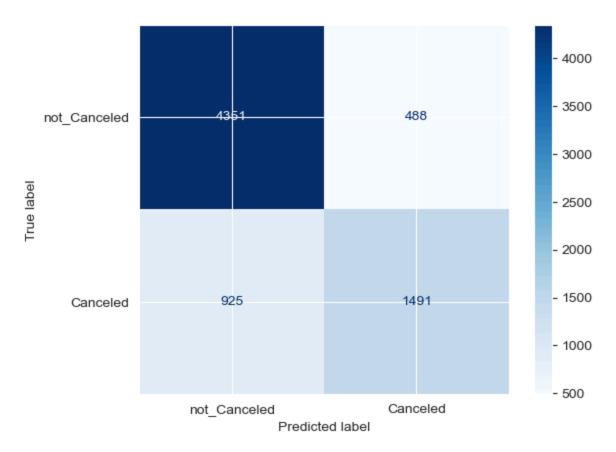
```
plt.title("Logistic Regression Decision Boundary with PCA")
plt.xlabel("PCA Component 1")
plt.ylabel("PCA Component 2")
plt.show()
```

Logistic Regression Decision Boundary with PCA



```
In [41]: # Confusion Matrix
cm = confusion_matrix(y_test, predict)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['not_Canceled',
disp.plot(cmap='Blues')
```

Out[41]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x179f06d7d10>



```
In [42]: # Accuracy
    accuracy = accuracy_score(y_test,predict)
    print(f"Model Accuracy: {accuracy*100:.2f}%")
```

Model Accuracy: 80.52%